

TWODAY MASTERCLASS

Design and development of autonomous AI-agents



Thomas Martinsen
Global Principal Advisor
Global Digital Engineering & Global AI



Lars Anker Daater
Senior Software Architect
Data & AI

Pan-Nordic footprint

EMPLOYEES

3000

LOCATIONS

34



Agenda

Design and development of **autonomous AI-agents**

Each topic takes 45-75 min

Theory and demo

20-30 min
lead by Thomas

Hands-on exercise

20-60 min
lead by Lars

Break are 15 min

Lunch is 30 min

Slides, demos, and exercises

<https://bit.ly/4jHw9fb>

- Introduction to the masterclass
- break 
- Agentic engineering
- Adapting knowledge, skills, and agents for engineering
- lunch 
- Platform for building and operating agents
- Agent development fundamentals
- break 
- Advanced agent design and orchestration
- Putting the pieces together
- Wrap up

Introduction


Evolution of AI

Traditional AI

- Rule-based Systems
- Symbolic AI
- Machine Learning
- Deep Learning
- Regression
- Classification
- Forecasting



Modern AI



- Generative AI
- Agentic AI
- ...




Strong AI

- AGI
- Super Intelligence
- Singularity

Relevant for all companies
and all employees

We are in the shift from manual workflows to automation and having AI at the core



Agentic impact on an organization



From workflows to goals

Instead of defining every step, organizations define intent, constraints, policies, and success criteria — agents determine parts of the execution path.



From systems of record to systems of action

Systems no longer only store and expose information; agents interpret context, make decisions, and initiate actions across systems.



From human execution to human oversight

Humans increasingly define boundaries, supervise outcomes, handle exceptions, and intervene where risk requires it.



From deterministic to probabilistic systems

We move from systems where the same input normally produces the same output to systems whose behavior must be continuously evaluated, constrained, and monitored.



Operating model

Authority moves closer to execution. Organizations must define what agents may decide, what requires approval, and when humans must intervene.



People & skills

Roles shift from executing tasks toward defining intent, supervising agents, evaluating outcomes, and managing exceptions.



Technology & data


APIs, identity, permissions, data quality, observability, and integration become prerequisites for safe autonomy — legacy systems increasingly limit what agents can do.



Governance & accountability


Organizations must be able to explain who authorized an agent, what it was allowed to do, what information it used, what actions it took, and who remains accountable.

Agents are the mechanism – not the destination




Agentic engineering

How development changes



Evolution in how we engineer AI systems



The risks of using AI in development

So much for 'trust but verify': Nearly half of software developers don't check AI-generated code – and 38% say it's because it takes longer than reviewing code produced by colleagues

A concerning number of developers are failing to check AI-generated code, exposing enterprises to huge security threats

Using AI might actually slow down experienced devs

News By Craig Hale published July 11, 2025

Experienced developers aren't getting any benefit from AI, report claims

The Register

AI coding tools make developers slower but they think they're faster, study finds

Predicted a 24% boost, but clocked a 19% drag

 [Thomas Claburn](#)

Fri 11 Jul 2025 // 22:41 UTC


Developers Lean on AI More, But Report Growing Doubts About Accuracy, Stack Overflow Survey Says

The Risks of Using AI in Software Development


Why AI Coding Tools Make Experienced Developers 19% Slower and How to Fix It

Oct 3, 2025  Molisha Shah

Evolution in how we engineer AI systems




Evolution in how we engineer AI systems





Agentic engineering

demo

Agent sessions



AI tools and resources



* Business (300 PRPM) or Enterprise (1500 PRPM)

Agentic engineering


exercise

Adapting knowledge, skills, and
agents for engineering


Deliver using agentic engineering

Structured, AI-orchestrated processes where the developer defines intent, specifications guide planning, agents execute tasks, and human validation gates ensure control and quality across the end-to-end flow.


Designed to be predictable, reusable, team-scalable, and governed.



Onboarding a new **member** to the team



Onboarding a new **agent** to the team



Structuring knowledge and primitives


demo

Adapting knowledge,
skills, and agents for
engineering

exercise


Platform for building and
operating agents

Agentic architecture




Multi agent orchestration


Single Agent




Network




Reflection




Supervisor




Hierarchical



Custom



Microsoft has different ways of building agents



Microsoft Foundry — agent platform

Model Catalog

GPT, Claude, Grok, Mistral, DeepSeek, NVIDIA Nemotron + fine-tuning for all

Foundry Agent Service

GA: agents that reason, plan & act across tools, data & workflows

Foundry Control Plane

End-to-end observability into agent behaviour — production governance

Foundry IQ

Pre-built RAG knowledge bases: SharePoint, Fabric, web — no custom pipeline

MCP Server (Cloud)

Hosted at mcp.ai.azure.com — connect from VS Code, zero local setup, Entra auth

Agent Memory

Preview: managed long-term memory across sessions — auto extract & retrieve

Deploy models in Microsoft Foundry

demo

More than just Azure / OpenAI models



OpenAI
(ChatGPT)



Google
(Gemini)



Anthropic
(Claude)



Mistral



X AI
(Grok)



Azure
OpenAI



Microsoft
Foundry



Amazon
Bedrock



Ollama
(Offline)



Foundry
Local



Open
Router



Together.
AI



Cohere




Hugging
Face



GitHub
Models

+11K

Model leaderboard



Compare models

←

Compare models

+ Add model

claude-opus-5-5 (v: 2)

×

gpt-6.1-sol (v: 2026-09-29)

×

grok-4.7 (v: 1)

×

Deploy

View details

Deploy

View details

Deploy

View details

Benchmarks	
Quality	-
Safety	-
Estimated cost	-
Throughput	-

Benchmarks	
Quality	0.85
Safety	3.33
Estimated cost	51.92
Throughput	65

Benchmarks	
Quality	-
Safety	-
Estimated cost	-
Throughput	-

Input	text,image,code
Output	text,image,code

Input	text,image
Output	text

Input	text,image
Output	text

Context	
Window	1.000.000 tokens
Max output	128.000 tokens
Training date	June 2026

Context	
Window	1.050.000 tokens
Max output	128.000 tokens
Training date	September 2026

Context	
Window	500.000 tokens
Max output	500.000 tokens
Training date	August 2026

Model pricing

GPT-6.1 Series

Model	Pricing (1M Tokens)	Priority Processing (1M Tokens)
GPT-6.1 Sol (short context) Global	Input: \$2 Cached Input: \$0.10 Cache writes: \$2.50 Output: \$10	N/A
GPT-6.1 Sol (short context) Data Zone	Input: \$2.20 Cached Input: \$0.11 Cache writes: \$2.75 Output: \$11	N/A
GPT-6.1 Sol (long context) Global	Input: \$4 Cached Input: \$0.20 Cache writes: \$5 Output: \$15	N/A
GPT-6.1 Sol (long context) Data Zone	Input: \$4.40 Cached Input: \$0.22 Cache writes: \$5.50 Output: \$16.50	N/A

Design and build agents in Microsoft Foundry

demo

Instruction in production

You help investigate workshop cases. Answer using only facts supplied in the conversation. If a company policy is requested but has not been supplied, say you do not have that policy. Never imply that you have read a case system or policy repository.

Company website discovery agent

Role

You are a company website discovery agent.

Your responsibility is to analyze a company's official website based on a supplied URL and return a concise, structured overview of the company together with the locations of the most important sections of the website.

Your task is limited to discovering:

- Basic company information
- The company's About page
- The company's News page
- The company's Careers page
- The company's general contact information

You must ****not**** collect news articles, job postings, employee information, or any other detailed content.

Input

The agent receives a single input:

Agent 1 – Find company details

get-company-details

PlaygroundDetailsTracesInsightsPreviewMonitorEvaluationOptimize

Model: gpt-6-lunaGlobal Standard deployment

Voice mode
Switch the agent to a voice-first experience.

Instructions
****Main products or services:****

- Autonomous warehouse robots
- Fleet management software
- Installation and support

Improve instructions

Tools
Build a unified endpoint for invoking tools and agents, leverage tool search to save input tokens. [Learn more about toolbox](#)

Web search

AddUpload files

Knowledge

MemoryPreview

Optimize

ChatYAML

9106t⌵⌵⌵Metrics⚙️

twoday.com

Twoday

Twoday is a technology company that helps public- and private-sector organizations transform through applied AI and advanced engineering.

Main products or services:

- Data and AI services, including analytics and automation
- Cloud platforms and security
- Custom software engineering and business applications

(twoday.com) 1

1 Who we are - twoday

gpt-6-luna6s9106t⌵Web searchmessage⌵⌵Traces👍👎📄🗑️

Message the agent...

🔄➡️

AI-generated content may be incorrect

Agent 2 – Find company information

← get-company-information ✓

Version: 3 (Today 10.29 pm) Save </> Continue in code Publish ▾

Playground Details Traces Insights Preview Monitor Evaluation Optimize

Model: gpt-6-sol Global Standard deployment ▾

Voice mode
Switch the agent to a voice-first experience. ☐

Instructions

Company website discovery agent
Role
You are a company website discovery agent.
Improve instructions

Tools

Build a unified endpoint for invoking tools and agents, leverage tool search to save input tokens. [Learn more about toolbox](#)

Web search ⓘ

Add ▾ Upload files

Knowledge

Memory Preview

Optimize ▾

Chat YAML

18131t ⓘ Metrics ▾ ⚙️ +

twoday.com

JSON Copy

```
1 {
2   "name": "Twoday",
3   "description": "Twoday provides AI, data, cloud, and software engineering services to public- and private-sect
4   "productsAndServices": "Applied AI, data services, cloud services, software engineering, and business applicat
5   "industriesOrMarkets": "Public and private sectors, primarily in the Nordics",
6   "websiteUrl": "https://www.twoday.com/",
7   "language": "English",
8   "region": "Nordics and Lithuania",
9   "contactUrl": "https://www.twoday.com/contact",
10  "careerUrl": "https://www.twoday.com/careers",
11  "newsUrl": "https://www.twoday.com/newsroom",
12  "aboutUrl": "https://www.twoday.com/about",
13  "generalEmails": [
14    "info.fi@twoday.com"
15  ],
16  "generalPhones": [
17    "+45 87 34 44 54"
18  ],
19  "analysisStatus": "completed",
20  "analyzedAt": "2026-10-06T20:29:53Z",
21  "sources": [
22    "https://www.twoday.com/",
```

Message the agent...
🔗 ↻ ➤

AI-generated content may be incorrect

Agent tools

Tools

Build a unified endpoint for your AI application that can handle multiple input tokens. [Learn more](#)

 Web search 

Web search uses Grounding with Bing, which has additional [costs](#) and terms: [terms of use](#) and [privacy statement](#). Customer data will flow outside the Azure compliance boundary. Learn more [here](#).

Add Model Context Protocol tool



Third-party services you connect to via Model Context Protocol are Non-Microsoft Products under the Product Terms. Use third-party services at your own risk and subject to third-party license terms and privacy policies.

Name * 

Provide a unique name

Remote MCP Server endpoint * 

Provide the remote MCP server endpoint

Authentication * 

Key-based



Credential * 

E.g. "Authorization" or "X"

:

E.g. "Bearer sk-abc123xyz" or "your-api-key-here"



+ Add key value pair

Connect

Cancel

Foundry Toolkit for VS Code



Foundry Toolkit for VS Code

Microsoft  microsoft.com |  1,495,066 |      (19)

Build AI agents and workflows in Microsoft Foundry, experiment with open or proprietary models.

☒ Auto Update 


What is Microsoft Foundry Toolkit

Microsoft Foundry Toolkit is the new name for AI Toolkit. It is a Visual Studio Code extension that helps you build AI agents quickly, with built-in access to [Microsoft Foundry](#) resources for model deployment, agent management, and more without leaving VS Code.


 **Note:** We have renamed this product to reflect its full integration with Microsoft Foundry. The previously separate Microsoft Foundry extension is now merged into this single extension. All existing features remain unchanged.

With Microsoft Foundry Toolkit you can:

-  **Discover and evaluate models** from providers including Microsoft Foundry, Foundry Local, Anthropic, OpenAI, GitHub, Google, and NVIDIA NIM, or run models locally with ONNX and Ollama.
-  **Build, test, and deploy AI agents** with Agent Builder for no-code prompt agents, or create code-based hosted agents with debugging, streaming visualization, and MCP tool integrations.
-  **Manage Microsoft Foundry resources** directly from VS Code: browse projects, deploy models from the catalog, and create, deploy, and test agents with the Microsoft Foundry Agent Service.



Using Microsoft Foundry as model provider



Model providers

Configured providers

Access models from other providers with your own API keys. Keys stay private and go straight to the provider.

+ Add provider



GitHub Copilot

From your GitHub Copilot subscription.

Provided by your GitHub account @thomas-martinsen_tw

Search providers...



Custom endpoint

Any OpenAI-compatible HTTP endpoint (vLLM, OpenRouter, fine-tune, etc.).



Anthropic

Hosted Claude models over the Messages API.



Azure OpenAI

Service deployments via your resource host, API version, and per-model deployment names.



Foundry Local

Local runtime with no API key, over an OpenAI-compatible endpoint.



Microsoft Foundry

OpenAI-compatible endpoints. Defaults to the Responses API for GPT-5.

exercise

Agent development fundamentals



Microsoft Agent Framework

Open-source engine for building and orchestrating intelligent AI agents

Open Standards &
Interoperability

Pipeline for
Research

Community-Driven &
Extensible by Design


Ready for
Production

aka.ms/AgentFramework

What is Microsoft Agent Framework


Standardized way to build agents, including tools and orchestration, reducing need for custom plumbing.

MICROSOFT AGENT FRAMEWORK




 OpenAI
  Gemini
  AI
  LLaMA
 + any LLM supporting IChatClient

Main concepts



Main concepts



Chat is more than just a chatbot

Assistant with human interaction
where the human is the initiator

*Can you summarize this report for me
and give me the highlights*

Autonomously in backend system
where a timer trigger can be the
initiator

*Every day @ 08:00, Search for AI news,
summarize it and send me an email*

Triggered by external system
where a webhook is the initiator

*When GitHub pull request is submitted,
make an AI review*

So "chat" is kind of wrong name – it is more "Request / Response"


Building chat agents using Microsoft Agent Framework

demo


Building agents using Microsoft Agent Framework

Getting ready

- Get endpoint and API key from project in Microsoft Foundry.
- Deploy the LLM models in Microsoft Foundry you want the agent to use.
- Store endpoint and model in configuration file, and ApiKey in secrets.json file.
- Add reference to relevant nuget-packages in project.



Endpoint: `https://<name>.services.ai.azure.com[/the rest is not necessary]`



A simple chat agent

```
var endpoint = Environment.GetEnvironmentVariable("FOUNDRY_PROJECT_ENDPOINT")
    ?? throw new InvalidOperationException("Set FOUNDRY_PROJECT_ENDPOINT.");
var model = Environment.GetEnvironmentVariable("FOUNDRY_MODEL")
    ?? throw new InvalidOperationException("Set FOUNDRY_MODEL.");
```

```
AIAgent agent = new AIProjectClient(new Uri(endpoint), new DefaultAzureCredential())
    .AsAIAgent(
        model: model,
        name: "CaseChat",
        instructions: """
        You help investigate workshop cases. Answer using only facts supplied in the conversation.
        If a company policy is requested but has not been supplied, say you do not have that policy.
        Never imply that you have read a case system or policy repository.
        """);
```

system prompt

AgentSession session = await agent.CreateSessionAsync(); ← session for conversation history

```
Console.WriteLine("Question 1: general case discussion");
```

```
Console.WriteLine(await agent.RunAsync(
```

user prompt

```
"A customer says their laptop will not start and they need a replacement tomorrow. What should I ask next?",
session));
```

```
Console.WriteLine("\nQuestion 2: information the chat cannot know yet");
```

```
Console.WriteLine(await agent.RunAsync(
    "What does our internal replacement policy require for this customer?", session));
```

Token cost

```
var input = usage.InputTokenCount;  
var output = usage.OutputTokenCount;  
var total = usage.TotalTokenCount ?? (input.HasValue && output.HasValue ? input + output : null);  
Console.WriteLine($"Token usage: In = {Format(input)} | Out = {Format(output)} | Total = {Format(total)}");
```

Interactive case chat. Press Enter to send a message; press Esc to exit.

You: hello

Agent: Hello! How can I help with your workshop case?

Token usage: In = 59 | Out = 15 | Total = 74

You: hello

Agent: Hi! What can I help you with?

Token usage: In = 81 | Out = 13 | Total = 94

Microsoft Agent Framework engineering skills

Microsoft Agent Framework Integration Instructions

Use these instructions when implementing a Microsoft Agent Framework integration backed by Azure OpenAI in a .NET backend. This document is intended to be portable across projects, while still reflecting the integration shape used here.

Goal

Build the integration as a thin, reusable backend service that:

- owns all Microsoft Agent Framework and Azure OpenAI SDK usage in one project or module
- exposes one non-streaming execution method and one streaming execution method
- is registered through dependency injection
- is configured from application settings or secret-backed environment variables
- can be consumed by API endpoints, background jobs, console hosts, or other backend entry points without duplicating SDK setup

Required architecture

- Create a dedicated backend project or module for agent integration. Do not scatter SDK calls across controllers, endpoints, or business services.
- Keep the public surface small:
 - one service interface
 - one request contract
 - one options or configuration contract
 - one concrete Azure-backed implementation
 - one DI registration helper
- Treat HTTP endpoints as adapters only. They should translate transport contracts into the shared agent request contract and delegate to the service.


Creating an engineering skill for Agent Framework

demo


exercise

Advanced agent design and orchestration

Main concepts



Two types of tools



Simple tool call example

```
var cases = new CaseLookup();
AIAgent agent = new AIProjectClient(new Uri(endpoint), new DefaultAzureCredential())
    .AsAIAgent(
        model: model,
        name: "CaseWithFunctionTool",
        instructions: """
            You investigate workshop cases but never update them. Call GetCase for the requested case ID.
            Use only the returned facts. If the tool fails, explain the failure without inventing a case.
            """,
        tools: [AIFunctionFactory.Create(cases.GetCase)]);
```

available tools which the agent can decide to call or not

1 reference

```
public sealed class CaseLookup
{
    [Description("Read one workshop case by its case ID. This tool never updates the case.")]
    1 reference
    public string GetCase([Description("Case ID, for example C-100.") string caseId)
    {
        Console.WriteLine($"GetCase called with {caseId}");
        return caseId switch
        {
            "C-100" => "C-100: Customer laptop will not start. Contact number 555-0100. Replacement is needed tomorrow.",
            "C-101" => "C-101: Employee cannot access portal. Affected account is unknown.",
            _ => throw new ArgumentException("Unknown case ID", nameof(caseId)),
        };
    }
}
```

Log tool calls via middleware

```
static async ValueTask<object?> LogToolCallAsync(
    AIAgent _,
    FunctionInvocationContext context,
    Func<FunctionInvocationContext, CancellationToken, ValueTask<object?>> next,
    CancellationToken cancellationToken)
{
    string arguments = context.Arguments.Count == 0
        ? ""
        : " (Args: " + string.Join(", ", context.Arguments.Select(argument =>
            $"{argument.Key} = {argument.Value}")) + ")";
    Console.WriteLine($"Tool call: '{context.Function.Name}'{arguments}");
    return await next(context, cancellationToken);
}
```

Internal C# function tool:

Tool call: 'GetCase' (Args: caseId = C-100)


Ask: "What happens when you press the power button?"

Token usage: In = 291 | Out = 87 | Total = 378

Tool calling

demo

Main concepts



Why do we need structured output

```
> What is the top 10 movies according to IMDB?
```

Certainly! As of **June 2024**, here are the **top 10 movies on IMDb** along with their **directors**:

1. **The Shawshank Redemption (1994)**
Director: Frank Darabont

2. **The Godfather (1972)**
Director: Francis Ford Coppola

3. **The Dark Knight (2008)**
Director: Christopher Nolan

4. **The Godfather Part II (1974)**
Director: Francis Ford Coppola

5. **12 Angry Men (1957)**
Director: Sidney Lumet

6. **Schindler's List (1993)**
Director: Steven Spielberg

7. **The Lord of the Rings: The Return of the King (2003)**
Director: Peter Jackson

8. **Pulp Fiction (1994)**
Director: Quentin Tarantino

9. **The Lord of the Rings: The Fellowship of the Ring (2001)**
Director: Peter Jackson

10. **The Good, the Bad and the Ugly (1966)**
Director: Sergio Leone

As of **June 2024**, here are the **top 10 movies according to IMDb's user ratings** (IMDb Top 250), though ranks can fluctuate slightly over time:

1. **The Shawshank Redemption (1994)** - 9.3
2. **The Godfather (1972)** - 9.2
3. **The Dark Knight (2008)** - 9.0
4. **The Godfather Part II (1974)** - 9.0
5. **12 Angry Men (1957)** - 8.9
6. **Schindler's List (1993)** - 8.9
7. **The Lord of the Rings: The Return of the King (2003)** - 8.9
8. **Pulp Fiction (1994)** - 8.8
9. **The Lord of the Rings: The Fellowship of the Ring (2001)** - 8.8
10. **The Good, the Bad and the Ugly (1966)** - 8.8

Note: Positions are subject to slight changes, but these titles have held firm at the top for some time. For the absolute latest, please refer to the [IMDb Top 250](<https://www.imdb.com/chart/top>).

For the most up-to-date list and details, always check the [IMDb Top 250](<https://www.imdb.com/chart/top>) directly.

Enabling structured output

```
AIAgent agent = new AIProjectClient(new Uri(endpoint), new DefaultAzureCredential())
    .AsAIAgent(model: model, name: "CaseClassifier", instructions: ""
        Call GetCase for the requested case ID, then return a structured triage result from those facts.
        Do not invent a policy or claim that you have updated a case.
        Use NeedsMoreInformation when facts are insufficient for a recommendation.
        """, tools: [AIFunctionFactory.Create(new CaseLookup().GetCase)]);
```

```
AgentResponse<CaseTriage> response = await agent.RunAsync<CaseTriage>(
    "Read case C-100 and return structured triage.");
```

```
TokenUsageReporter.Write(response);
```

```
CaseTriage result = response.Result;
```

```
if (result.CaseId != "C-100" || string.IsNullOrWhiteSpace(result.Summary)
    throw new InvalidDataException("The structured output lacks the ex
```

3 references

```
public sealed class CaseTriage
```

```
{
```

1 reference

```
public string CaseId { get; set; } = "";
```

1 reference

```
public string Summary { get; set; } = "";
```

0 references

```
public string Priority { get; set; } = "";
```

0 references

```
public bool NeedsMoreInformation { get; set; }
```

0 references


```
public string NextQuestion { get; set; } = "";
```

```
}
```


Structured output

demo


Main concepts




Retrival Augmented Generation



LLM grounded data





Adding knowledge



Limitations in RAG


Context Window Max Size:

128,000 – 1,000,000 Tokens (~ 500-2000 pages of text)
(various from LLM to LLM)



...

Vector search



RAG

demo

exercise


Putting the pieces together

UI for AI

Assistants and agents

AI engines

LLMs



Calling Foundry agents from code

Call get-company-information in code



Send a message to this agent and get a response.

Agent name

Project endpoint

get-company-information



https://niq-shared-sweden-resource.servic



Call this agent

Language C# ▾



```

1 // dotnet add package Azure.AI.Projects --version 2.0.0-beta.2
2 using Azure.AI.Projects;
3 using Azure.AI.Projects.Agents;
4 using Azure.AI.Extensions.OpenAI;
5 using Azure.Identity;
6 using OpenAI.Responses;
7
8 #pragma warning disable OPENAI001
9
10 const string endpoint = "https://niq-shared-sweden-resource.services.ai.azure.com/api/projects/niq-shared-sweden";
11 const string agentName = "get-company-information";
12 const string agentVersion = "6";
13
14 // Connect to your project using the endpoint from your project page
15 // The AzureCliCredential will use your logged-in Azure CLI identity, make sure to run `az login` first
16 AIProjectClient projectClient = new(endpoint: new Uri(endpoint), tokenProvider: new DefaultAzureCredential());
17
18 AgentReference agentReference = new(name: agentName, version: agentVersion);
19 ProjectResponsesClient responseClient = projectClient.OpenAI.GetProjectResponsesClientForAgent(agentReference);
20 // Use the agent to generate a response
21 ResponseResult response = responseClient.CreateResponse(
22     "Hello! Tell me a joke."
23 );
24
25 Console.WriteLine(response.GetOutputText());

```

demo

exercise

Masterclass wrap up

TWODAY MASTERCLASS

Questions

Engineering autonomous AI-agents

You're always welcome to reach out with **questions** or **comments** after this masterclass.



Thomas Martinsen
Global Principal Advisor
thomas.martinsen@twoday.com



Lars Anker Daater
Senior Software Architect
lars.daater@twoday.com

twoday